

Basic Statistics in R

ECO 6416

2022-08-28

Contents

1	Univariate Analysis	2
1.1	Center	2
1.2	Shape	3
1.3	Spread	4
1.4	Easier Way	6
2	Bivariate Analysis	7
2.1	Correlation	7
2.2	Scatterplots	7
2.3	Time series graphs	8
3	More Visual Displays!	9
3.1	Data Tables	9
3.2	Bar Charts	10
3.3	Stem and Leaf Plots	10
4	Fancier Output	11
4.1	Table Output	11
4.2	Summary Statistics	12
4.3	Histograms	12
4.4	Boxplots	13
4.5	Correlation	13
4.6	Scatterplots	14
4.7	Trends	15

Here are all the packages needed to get started.

```
library(gt)
library(tidyverse)
library(gtsummary)
library(plotly)
library(readxl)
library(plotly)
library(corrplot)
```

```
sessionInfo()
```

```
## R version 4.2.1 (2022-06-23 ucrt)
## Platform: x86_64-w64-mingw32/x64 (64-bit)
```

```
## Running under: Windows 10 x64 (build 19044)
##
## Matrix products: default
##
## locale:
## [1] LC_COLLATE=English_United States.utf8
## [2] LC_CTYPE=English_United States.utf8
## [3] LC_MONETARY=English_United States.utf8
## [4] LC_NUMERIC=C
## [5] LC_TIME=English_United States.utf8
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] corrplot_0.92   readxl_1.4.1    plotly_4.10.0   gtsummary_1.6.1
## [5] forcats_0.5.2   stringr_1.4.0    dplyr_1.0.9      purrr_0.3.4
## [9] readr_2.1.2     tidyr_1.2.0     tibble_3.1.7    ggplot2_3.3.6
## [13] tidyverse_1.3.2 gt_0.7.0
##
## loaded via a namespace (and not attached):
## [1] lubridate_1.8.0   assertthat_0.2.1  digest_0.6.29
## [4] utf8_1.2.2        R6_2.5.1          cellranger_1.1.0
## [7] backports_1.4.1   reprex_2.0.2      evaluate_0.15
## [10] httr_1.4.3        pillar_1.7.0      rlang_1.0.3
## [13] lazyeval_0.2.2    googlesheets4_1.0.1 rstudioapi_0.14
## [16] data.table_1.14.2  rmarkdown_2.14    googledrive_2.0.0
## [19] htmlwidgets_1.5.4  munsell_0.5.0     broom_1.0.0
## [22] compiler_4.2.1    modelr_0.1.9      xfun_0.31
## [25] pkgconfig_2.0.3    htmltools_0.5.2   tidyselect_1.1.2
## [28] viridisLite_0.4.0  fansi_1.0.3       crayon_1.5.1
## [31] tzdb_0.3.0        dbplyr_2.2.1      withr_2.5.0
## [34] grid_4.2.1        jsonlite_1.8.0    gtable_0.3.0
## [37] lifecycle_1.0.1   DBI_1.1.3         magrittr_2.0.3
## [40] scales_1.2.0      cli_3.3.0         stringi_1.7.8
## [43] broom.helpers_1.8.0 fs_1.5.2          xml2_1.3.3
## [46] ellipsis_0.3.2     generics_0.1.3    vctrs_0.4.1
## [49] tools_4.2.1       glue_1.6.2        hms_1.1.1
## [52] fastmap_1.1.0     yaml_2.3.5        colorspace_2.0-3
## [55] gargle_1.2.0      rvest_1.0.3       knitr_1.39
## [58] haven_2.5.1
```

1 Univariate Analysis

In univariate analysis, we look at a single variable and describe 3 different things:

- Center
- Shape
- Spread

1.1 Center

Helps explain where the middle of the data is. This can be measured in 3 main ways.

1.1.1 Mean

```
grades <- c(78,79,80,81,82)
mean(grades)
```

```
## [1] 80
```

1.1.2 Median

```
median(grades)
```

```
## [1] 80
```

1.1.3 Mode

There isn't an easy way of doing this, so I created a function instead.

```
getModes <- function(x) {
  ux <- unique(x)
  tab <- tabulate(match(x, ux))
  ux[tab == max(tab)]
}
```

```
getModes(grades)
```


```
## [1] 78 79 80 81 82
```

Since there is no value that occurs most frequently, they all show. Mode is rarely used.

1.2 Shape

For this, I am going to randomly generate 100 exam scores.

```
grades <- rnorm(100,mean = 75, sd = 5 )
hist(grades)
```



You can mess with the `breaks =` argument to get different numbers of bins.

1.3 Spread

Let's look at a different dataset. Starwars character heights.

```
var(starwars$height)
```

```
## [1] NA
```

What happened? We have nulls in our data, so we cannot calculate the variance until we tell the system to ignore null values

```
var(starwars$height, na.rm = TRUE)
```

```
## [1] 1208.983
```

```
sd(starwars$height, na.rm = TRUE)
```

```
## [1] 34.77043
```

```
IQR(starwars$height, na.rm = TRUE)
```

```
## [1] 24
```


```
range(starwars$height, na.rm = TRUE)
```

```
## [1] 66 264
```

1.3.1 Visualizing Spread

You can visualize the spread as well this with a boxplot.

```
boxplot(starwars$height)
```

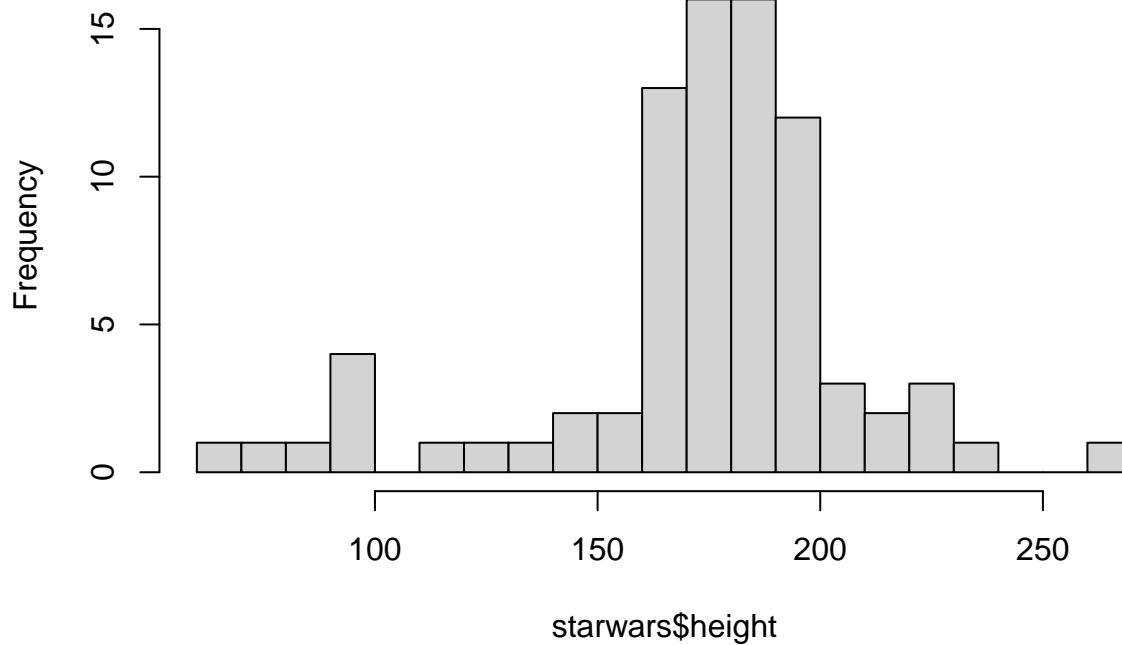


You can see that a lot of them fall outside of the fences.

You could also visualize this with a histogram

```
hist(starwars$height, breaks = "fd")
```

Histogram of starwars\$height



1.4 Easier Way

To get most of the items of interest, you can simply use the `summary()` function. The standard deviation nor the variance is displayed, so you would still need those.

```
summary(mtcars)
```

```
##      mpg          cyl        disp         hp
##  Min.   :10.40   Min.    :4.000   Min.    : 71.1   Min.    : 52.0
##  1st Qu.:15.43   1st Qu. :4.000   1st Qu. :120.8   1st Qu. : 96.5
##  Median :19.20   Median :6.000   Median :196.3   Median :123.0
##  Mean   :20.09   Mean    :6.188   Mean    :230.7   Mean    :146.7
##  3rd Qu.:22.80   3rd Qu. :8.000   3rd Qu. :326.0   3rd Qu. :180.0
##  Max.   :33.90   Max.    :8.000   Max.    :472.0   Max.    :335.0
##      drat          wt          qsec         vs
##  Min.   :2.760   Min.    :1.513   Min.    :14.50   Min.    :0.0000
##  1st Qu.:3.080   1st Qu. :2.581   1st Qu. :16.89   1st Qu. :0.0000
##  Median :3.695   Median :3.325   Median :17.71   Median :0.0000
##  Mean   :3.597   Mean    :3.217   Mean    :17.85   Mean    :0.4375
##  3rd Qu.:3.920   3rd Qu. :3.610   3rd Qu. :18.90   3rd Qu. :1.0000
##  Max.   :4.930   Max.    :5.424   Max.    :22.90   Max.    :1.0000
##      am          gear        carb
##  Min.   :0.0000   Min.    :3.000   Min.    :1.000
##  1st Qu.:0.0000   1st Qu. :3.000   1st Qu. :2.000
##  Median :0.0000   Median :4.000   Median :2.000
##  Mean   :0.4062   Mean    :3.688   Mean    :2.812
##  3rd Qu.:1.0000   3rd Qu. :4.000   3rd Qu. :4.000
```

```
## Max.      :1.0000   Max.      :5.000   Max.      :8.000
```

2 Bivariate Analysis

When looking at 2 variables, we use correlation, scatterplots, and time series graphs.

2.1 Correlation

```
cor(mtcars$mpg, mtcars$hp)
```


```
## [1] -0.7761684
```

These are strongly negatively correlated. We might say, horsepower has a inverse relationship with mpg, or mpg has a negative relationship with horsepower.

2.2 Scatterplots

We can also visualize this relationship with a scatterplot.


```
plot(mtcars$hp, mtcars$mpg)
```



we can also make this look nicer with some labels

```
plot(mtcars$hp, mtcars$mpg,  
     main = "Scatterplot of Horsepower and Miles Per Gallon",  
     xlab = "Horsepower",  
     ylab = "Miles Per Gallon")
```

Scatterplot of Horsepower and Miles Per Gallon




2.3 Time series graphs

Since the data we've seen thus far has been only cross section, let's bring in the data from the previous module. `Wealth1Percent.xlsx`

```
Wealth1percent <- read_excel("../Data/Wealth1percent.xlsx",  
                             col_types = c("date", "numeric", "numeric"))
```

To add the trend line, we need to do something that we haven't discussed yet. Don't worry, I'll explain it in the next module.

```
plot(Wealth1percent$quarter, Wealth1percent$Share, type = "l")  
abline(lm(Share ~ quarter, data = Wealth1percent), lty = 2)
```

3 More Visual Displays!

Here are some more ways to display data

3.1 Data Tables

We can represent data in categorical fashion:

```
table(starwars$hair_color)
```

```
##
##      auburn auburn, grey auburn, white      black      blond
##          1           1           1          13           3
##      blonde      brown  brown, grey      grey      none
##          1          18           1           1          37
##      unknown      white
##          1           4
```

Or quantitative

```
bins <- seq(10,34,by = 2)
```

```
mpg <- cut(mtcars$mpg,bins)
```


```
table(mpg)
```

```
## mpg
```

```
## (10,12] (12,14] (14,16] (16,18] (18,20] (20,22] (22,24] (24,26] (26,28] (28,30]
##      2      1      7      3      5      5      2      2      1      0
## (30,32] (32,34]
##      2      2
```

3.2 Bar Charts

```
n <- table(mtcars$gear)
barplot(n,xlab="Number of Gears")
```



3.3 Stem and Leaf Plots

We can add `scale = 3` to make this line up properly

```
stem(mtcars$hp, scale = 3)
```

```
##
## The decimal point is 1 digit(s) to the right of the |
##
## 5 | 2
## 6 | 2566
## 7 |
## 8 |
## 9 | 1357
## 10 | 59
## 11 | 0003
```

```
## 12 | 33
## 13 |
## 14 |
## 15 | 00
## 16 |
## 17 | 555
## 18 | 000
## 19 |
## 20 | 5
## 21 | 5
## 22 |
## 23 | 0
## 24 | 55
## 25 |
## 26 | 4
## 27 |
## 28 |
## 29 |
## 30 |
## 31 |
## 32 |
## 33 | 5
```

4 Fancier Output

Check out this fun stuff! Makes things look much cleaner.

4.1 Table Output

Categorical

```
hair <- table(starwars$hair_color)%>%
  data.frame()

colnames(hair)[1] <- "Hair Color"

gt(hair)
```

Hair Color	Freq
auburn	1
auburn, grey	1
auburn, white	1
black	13
blond	3
blonde	1
brown	18
brown, grey	1
grey	1
none	37
unknown	1
white	4

Or quantitative

```
bins <- seq(10,34,by = 2)

mpg <- cut(mtcars$mpg,bins)

table(mpg)%>%
  data.frame() %>%
  gt()
```

4.2 Summary Statistics

Using our classic `mtcars` dataset.

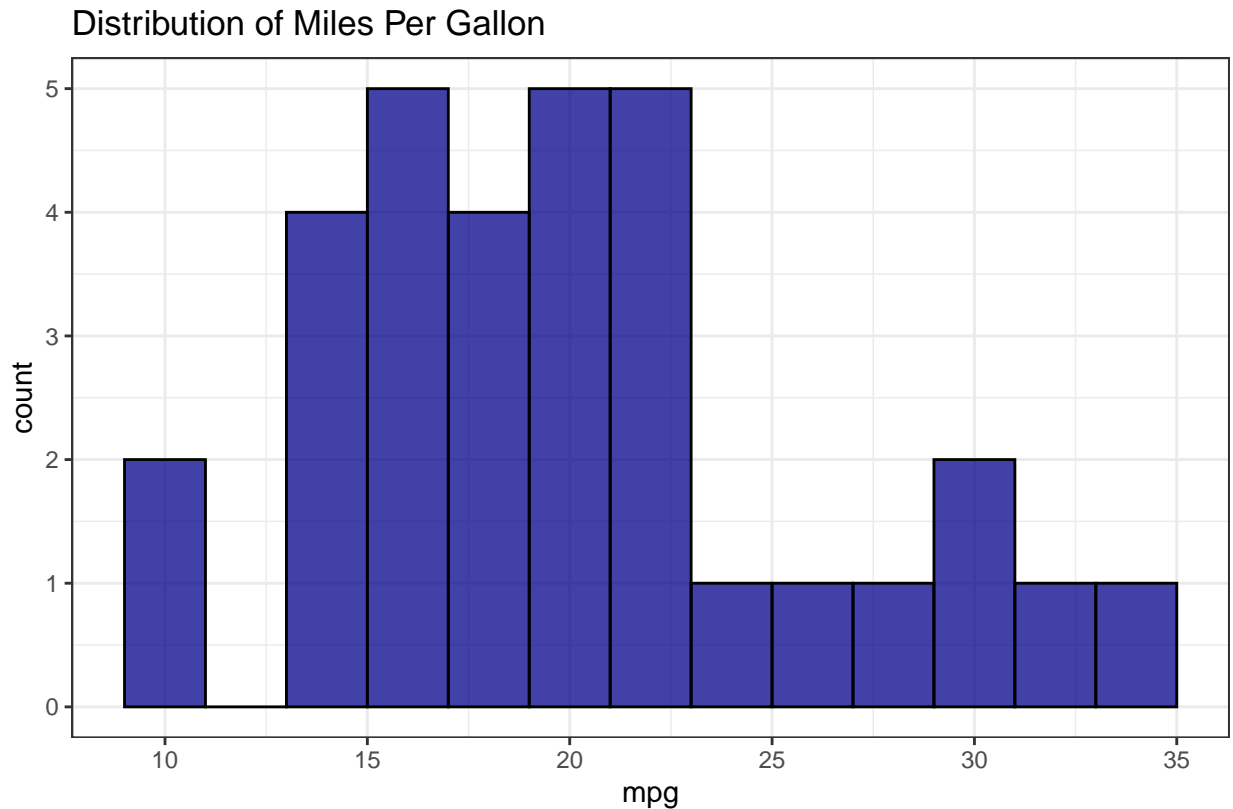
```
mtcars%>% select(mpg, cyl, hp) %>%
  tbl_summary(
    statistic = list(all_continuous() ~ c("{mean} ({sd})",
                                           "{median} ({p25}, {p75})",
                                           "{min}, {max}"),
                    all_categorical() ~ "{n} / {N} ({p}%)",
                    type = all_continuous() ~ "continuous2"
  )
```

Characteristic	N = 32
mpg	
Mean (SD)	20.1 (6.0)
Median (IQR)	19.2 (15.4, 22.8)
Range	10.4, 33.9
cyl	
4	11 / 32 (34%)
6	7 / 32 (22%)
8	14 / 32 (44%)
hp	
Mean (SD)	147 (69)
Median (IQR)	123 (96, 180)
Range	52, 335

4.3 Histograms

4.3.1 ggplot

```
ggplot(mtcars, aes(mpg))+
  geom_histogram(binwidth = 2,col = 'black', fill = 'darkblue', alpha = 0.75)+
  labs(title = 'Distribution of Miles Per Gallon', caption = "1974 Motor Trend US Magazine")+
  theme_bw()
```



1974 Motor Trend US Magazine

4.3.2 plot_ly

Since this is an interactive graph, it will not show up in a .pdf file, but it is great to look at in an .html document.

```
plot_ly(x = ~mtcars$mpg, type = "histogram", alpha = 0.6) %>%
  layout(title = 'Distribution of Miles Per Gallon',
         xaxis = list(title = 'Miles Per Gallon'),
         yaxis = list(title = 'Count'))
```

4.4 Boxplots

I think the plotly version of a boxplot is superior since it is interactive.

```
plot_ly(y = starwars$height, type = 'box', name = 'Height [cm]', text = starwars$name) %>%
  layout(title = 'Distribution of Star Wars Character Heights')
```

4.5 Correlation

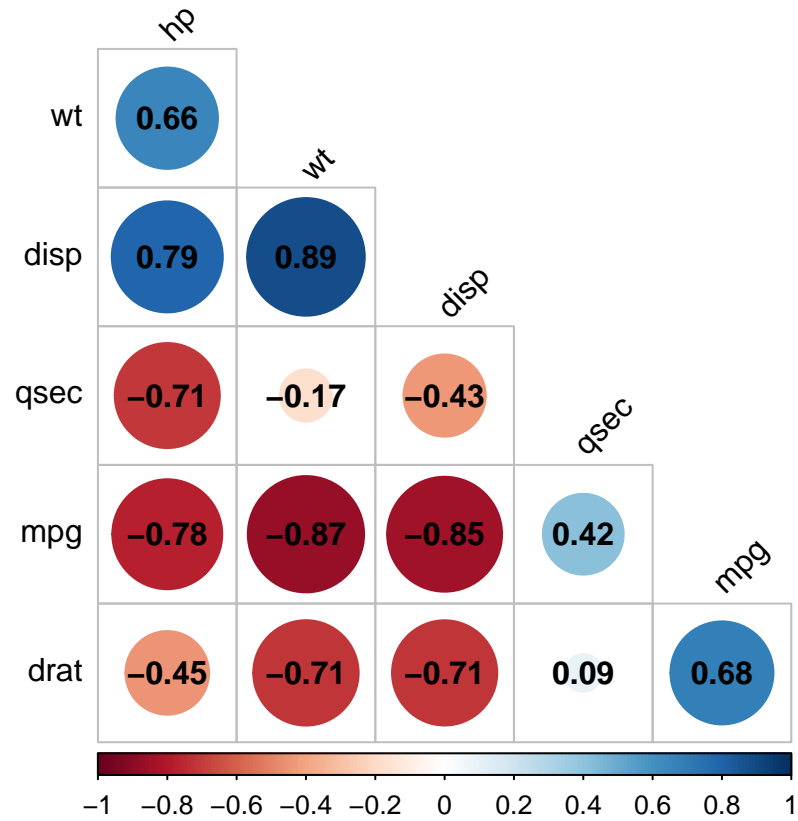
```
reduced <- mtcars %>%
  select(mpg, hp, wt, qsec, disp, drat)

corrplot(cor(reduced),
         type = "lower",
         order = "hclust",
         tl.col = "black",
```

```

tl.srt = 45,
addCoef.col = "black",
diag = FALSE)

```



4.6 Scatterplots


4.6.1 ggplot

```

ggplot(starwars,aes(height, mass))+
  geom_point(color = 'gray40')+
  theme_bw()+
  labs(title = "Relationship between Mass and Height of Star Wars Characters")

```

Relationship between Mass and Height of Star Wars Characters



4.6.2 plotly


```
plot_ly(starwars, y = ~mass, x = ~height, type = 'scatter', text = ~name, mode = "markers")
```

4.7 Trends

4.7.1 ggplot

```
ggplot(Wealth1percent, aes(quarter, Share))+  
  geom_line(color = 'gray40', alpha = 0.75)+  
  geom_smooth(method = "lm", se = F, color = 'darkblue', linetype = 'dashed')+  
  theme_bw()+  
  labs(title = "Share of Total Net Worth Held by the Top 1%",  
        subtitle = 'from 1989-2022',  
        x = "Date",  
        y = "Share")
```

Share of Total Net Worth Held by the Top 1%
from 1989–2022



4.7.2 plotly

To add the trend line, it gets quite tricky.

```
model <- lm(Share~quarter,Wealth1percent)
Trend <- predict(model,data = Wealth1percent$quarter)

plot_ly(Wealth1percent, x = ~quarter, y = ~Share, type = 'scatter', mode = 'lines', name = "Share") %>%
add_trace(y = ~Trend, name = 'Trend', mode = 'lines')
```